


```
[20]: #Put your code here
# parameters to search
params = {'C': [2 ** c for c in range(-14, 15)]}

# classifier
lr_clf = GridSearchCV(LogisticRegression(), params, cv=10)
lr_clf.fit(data_train, label_train)

# cross validated accuracy and best parameters
bestparams = lr_clf.best_params_
lr_crossvalacc = lr_clf.cv_results_['mean_test_score'][np.argmax(np.array(lr_clf.cv_results_['params']
))] == bestparams][0][0]

# print results
print("C: 2^{0}".format(int(np.log2(bestparams['C']))))
print("Cross Validation Error: {}".format(1 - lr_crossvalacc))

C: 2^0
Cross Validation Error: 0.02145505279034726
```

[C: 2⁰ (1) and Error: 0.02145 or 2.145 %]

Based on the classifiers you selected thusfar for Linear SVM, SVM + Gaussian RBF and Logistic Regression, which classifier would you pick? Make sure to take into account error, the application and computational considerations. **(5 points)**

[Based on the above classifiers I would pick the radial basis function (RBF) SVM because it appears to have the smallest cross-validation error. However, the above implementation performs lots of computations (for classifying as well as searching for the best parameters) but the gain in Cross Validation accuracy appears to significant]

Train the classifier selected above on the whole training set. Then, estimate the prediction error using the test set. What is your estimate of the prediction error? How does it compare to the cross-validation error? **(10 points)**

```
In [21]: #Put your code here
# classifier
clf = svm.SVC(kernel='rbf', C=2**3, gamma=2**(-3))

# fit SVC classifier on whole training set
clf.fit(data_train, label_train)

# prediction on test data
predtestlabels = clf.predict(data_test)

# print test error
print("Test Error: {}".format(np.mean(predtestlabels != label_test)))

Test Error: 0.017543859649122806

[Test Error: 0.01754 or 1.754 %]
```

Do you think the 0.1-loss is appropriate performance measure to report, in this case? If so, why? If not, how would you measure performance? **(5 points)**

[The 0.1-loss isn't an appropriate performance measure because in the case of SVM we are attempting to maximize the definition of the "margin" which is used to separate the data. Therefore, allowing for gradation (rather than a binary measure) is a better measure of performance because it indicates how far away our data points are from a decision boundry. A better measure of performance is the hinge loss function which controls the size of the margin]

And this concludes Lab 3! Congratulations!